

# AUTOTUNING HLS FOR FPGAS USING OPENTUNER AND LEGUP

---

Pedro Bruel ([phrb@ime.usp.br](mailto:phrb@ime.usp.br))

**Alfredo Goldman** ([gold@ime.usp.br](mailto:gold@ime.usp.br))

Sai Rahul Chalamalasetti ([gold@ime.usp.br](mailto:gold@ime.usp.br))

Dejan Milojicic ([gold@ime.usp.br](mailto:gold@ime.usp.br))

ReConFig, December 5, 2017



*Institute of Mathematics and Statistics  
University of São Paulo*



1. FPGAs, HLS & Autotuning
2. Background
3. Experiments & Results
4. Conclusion



The slides and all source code are hosted at [GitHub](#):

- [Code & Data](#): `github.com/phrb/legup-tuner`
- [Slides](#): `github.com/phrb/slides-reconfig-2017-autotuning`

## FPGAs:

- Logic Blocks and Interconnections
- Reconfigurable

## Tradeoff:

- Energy Efficiency and Performance
- Programmability



## Using FPGAs in Bing:

1,632 Servers with FPGAs Running Bing Page Ranking Service (~30,000 lines of C++)

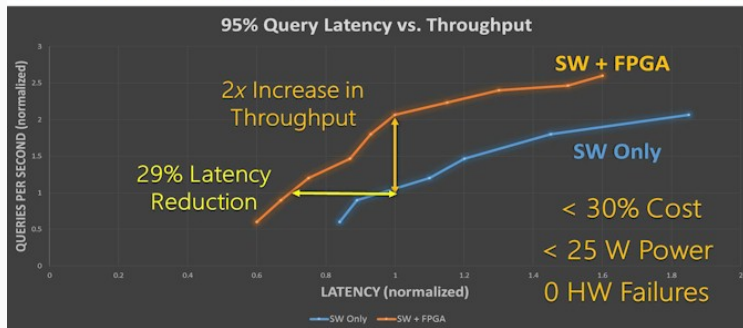


Image: [enterprisetech.com/2014/09/03/microsoft-using-fpgas-speed-bing-search/](http://enterprisetech.com/2014/09/03/microsoft-using-fpgas-speed-bing-search/) [Accessed in 27/11/17]

# FPGAs: HIGH-LEVEL SYNTHESIS

HLS can generate **lower-latency applications**:




Image: Karras, Kimon, Michaela Blott, and Kees Vissers. "High-Level Synthesis Case Study: Implementation of a Memcached Server." arXiv preprint arXiv:1408.5387 (2014).

# FPGAs: HIGH-LEVEL SYNTHESIS

Qualitatively, with less effort:

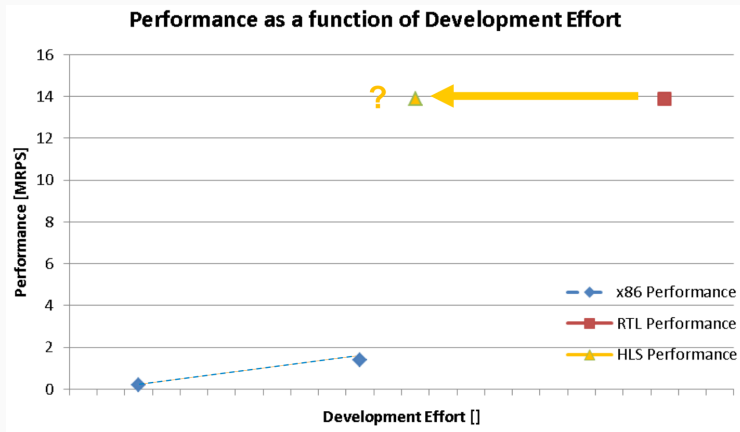


Image: Blott, Michaela, et al. "Achieving 10Gbps line-rate key-value stores with FPGAs." Presented as part of the 5th USENIX Workshop on Hot Topics in Cloud Computing. 2013.

This is an **old issue**:




Image: Smith, Steven W. "The scientist and engineer's guide to digital signal processing." 1997



# FPGAs: HIGH-LEVEL SYNTHESIS

## LegUp HLS flow:

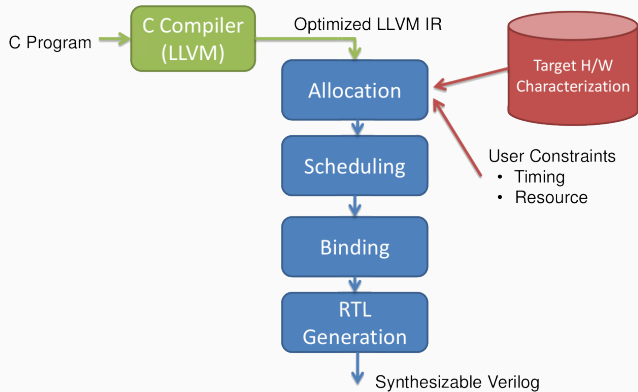



Image: Canis, Andrew Christopher. LegUp: Open-Source High-Level Synthesis Research Framework. Diss. University of Toronto, 2015.

## Why use autotuning for HLS?

- HLS aims to ease FPGA programming for software engineers
- But configuring HLS requires FPGA programming knowledge
- Autotuning can help software engineers use HLS


# FPGAs: AUTOTUNING FOR HIGH-LEVEL SYNTHESIS

Our HLS **compilation flow**:



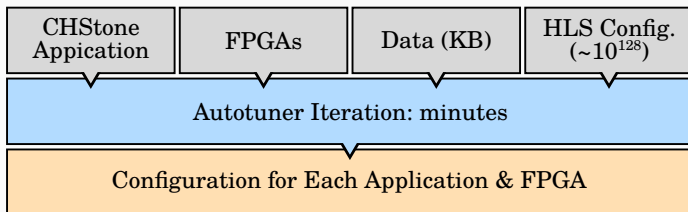
# FPGAs: AUTOTUNING FOR HIGH-LEVEL SYNTHESIS

Our **autotuner** setup:



# FPGAs: AUTOTUNING FOR HIGH-LEVEL SYNTHESIS

Our HLS **autotuning workflow**:



Huang, Qijing, *et al.*, "The effect of compiler optimizations on high-level synthesis-generated hardware." *ACM Transactions on Reconfigurable Technology and Systems* (2015):

- Optimizes parameters & ordering of a subset of LLVM passes
- Uses LegUp as an HLS tool
- Targets the Cyclone II device family
- Uses the CHStone benchmark
- Performs Exhaustive Search

*Xu et al., "A Parallel Bandit-Based Approach for Autotuning FPGA Compilation." FPGA (2017):*

- Does not optimize HLS
- Optimizes parameters of the Verilog-to-Routing (VTR) FPGA compilation tool
- Uses distributed OpenTuner instances

## C-based High-Level Synthesis [benchmark](#) (CHStone):

**Table 1:** Autotuned CHStone Applications

| Application | Short Description   |
|-------------|---|
| blowfish    | Symmetric-key block cypher                                |
| aes         | Advanced Encryption Algorithm (AES)                       |
| adpcm       | Adaptive Differential Pulse Code Modulation dec. and enc. |
| sha         | Secure Hash Algorithm (SHA)                               |
| motion      | Motion vector decoding from MPEG-2                        |
| mips        | Simplified MIPS processor                                 |
| gsm         | Predictive coding analysis of systems for mobile comms.   |
| dfsin       | Sine function for double-precision floating-point numbers |
| dfmul       | Double-precision floating-point multiplication            |
| dfdiv       | Double-precision floating-point division                  |
| dfadd       | Double-precision floating-point addition                  |



## Metric composition problem:

- 8 hardware metrics obtained from Quartus
- But the autotuner optimizes a single value

## Metric composition problem:

- 8 hardware metrics obtained from Quartus
- But the autotuner optimizes a single value

## Our solution so far:

- Normalize metrics using initial values
- Optimize the normalized sum of the 8 metrics

Normalized sum of metrics:

$$f(M, W) = \frac{\sum_{\substack{m_i \in M \\ w_i \in W}} w_i \left( \frac{m_i}{m_i^0} \right)}{\sum_{w_i \in W} w_i}$$

Where  $M$  is the set of hardware metrics and  $W$  is the set of weights

An **Optimization Scenario** consists of:

- An **optimization objective**: performance, area, . . .
- **Weights** for hardware metrics

An **Optimization Scenario** consists of:

- An **optimization objective**: performance, area, . . .
- **Weights** for hardware metrics

Our scenarios:

- 3 **specific** scenarios & 1 **balanced** scenario
- Weights: **powers of two** from 1: **irrelevant** to 8: **high**

# WEIGHTED OPTIMIZATION SCENARIOS

Table 2: **Weights** for **Optimization Scenarios** (High = 8, Medium = 4, Low = 2)

| Metric           | Area | Perf. & Lat | Performance | Balanced |
|------------------|------|-------------|-------------|----------|
| <i>LUT</i>       | High | Low         | Low         | Medium   |
| <i>Registers</i> | High | High        | Medium      | Medium   |
| <i>BRAMs</i>     | High | Low         | Low         | Medium   |
| <i>DSPs</i>      | High | Low         | Low         | Medium   |
| <i>FMax</i>      | Low  | High        | High        | Medium   |
| <i>Cycles</i>    | Low  | High        | Low         | Medium   |

## Experiments

- We performed 10 autotuning runs for each scenario
- Each run took 1.5 hours
- We kept track of how each hardware metric changed over time
- We used the Weighted Normalized Sum (WNS) as autotuning objective

## Presentation of results for each scenario:

- Mean of the relative changes, over 10 tuning runs
- For each hardware metric, and for WNS
- For each CHStone application
- Bluer is always better than Redder



## RESULTS: COMPARING DEFAULT & RANDOM STARTS

|                 |             |             |             |             |               |               |            |             |
|-----------------|-------------|-------------|-------------|-------------|---------------|---------------|------------|-------------|
| <i>aes</i>      | 0.79        | 0.91        | 1.00        | 0.56        | 1.00          | 0.47          | 1.00       | 1.12        |
| <i>adpcm</i>    | 0.68        | 1.13        | 1.00        | 0.54        | 1.00          | 0.56          | 0.60       | 0.98        |
| <i>sha</i>      | 1.00        | 1.03        | 1.00        | 0.82        | 1.00          | 0.55          | 1.00       | 0.89        |
| <i>motion</i>   | 1.02        | 1.00        | 0.60        | 0.85        | 1.00          | 0.57          | 1.00       | 0.94        |
| <i>mips</i>     | 1.00        | 0.93        | 1.00        | 0.44        | 1.00          | 0.45          | 0.98       | 1.24        |
| <i>gsm</i>      | 0.83        | 1.17        | 1.00        | 0.48        | 1.00          | 0.56          | 0.52       | 0.99        |
| <i>dfsint</i>   | 0.79        | 0.97        | 1.00        | 0.61        | 1.00          | 0.60          | 1.49       | 1.53        |
| <i>dfmul</i>    | 1.00        | 1.06        | 0.90        | 0.47        | 0.90          | 0.47          | 1.31       | 1.10        |
| <i>dfdiv</i>    | 0.83        | 1.07        | 0.80        | 0.73        | 0.80          | 0.65          | 1.32       | 1.49        |
| <i>dfadd</i>    | 1.00        | 0.94        | 1.00        | 0.82        | 1.00          | 0.71          | 1.00       | 0.93        |
| <i>blowfish</i> | —           | —           | —           | —           | —             | —             | —          | —           |
|                 | <i>LUTs</i> | <i>Pins</i> | <i>BRAM</i> | <i>Regs</i> | <i>Blocks</i> | <i>Cycles</i> | <i>DSP</i> | <i>FMax</i> |

**Figure 1:** Comparison of the absolute values for Random and Default starting points in the Balanced scenario

## RESULTS: THE BALANCED SCENARIO

|                 |            |             |             |             |             |               |               |            |             |
|-----------------|------------|-------------|-------------|-------------|-------------|---------------|---------------|------------|-------------|
| <i>aes</i>      | 0.94       | 0.90        | 1.00        | 1.00        | 0.97        | 1.00          | 0.98          | 1.00       | 0.76        |
| <i>adpcm</i>    | 0.84       | 0.73        | 1.13        | 1.00        | 0.91        | 1.00          | 1.05          | 0.63       | 0.49        |
| <i>sha</i>      | 0.98       | 1.00        | 1.03        | 1.00        | 0.96        | 1.00          | 0.86          | 1.00       | 0.97        |
| <i>motion</i>   | 0.98       | 0.97        | 1.00        | 1.00        | 0.99        | 1.00          | 0.95          | 1.00       | 0.95        |
| <i>mips</i>     | 0.95       | 1.00        | 1.07        | 1.00        | 0.90        | 1.00          | 1.01          | 0.80       | 0.89        |
| <i>gsm</i>      | 0.95       | 0.95        | 1.17        | 1.00        | 0.99        | 1.00          | 0.95          | 0.49       | 1.08        |
| <i>dfsint</i>   | 0.93       | 1.03        | 1.03        | 1.00        | 1.05        | 1.00          | 1.07          | 0.65       | 0.73        |
| <i>dfmul</i>    | 0.85       | 1.00        | 1.13        | 0.90        | 0.88        | 0.90          | 1.06          | 0.31       | 0.76        |
| <i>dfdiv</i>    | 0.84       | 1.00        | 1.07        | 0.80        | 1.15        | 0.80          | 1.34          | 0.41       | 0.57        |
| <i>dfadd</i>    | 1.00       | 1.00        | 1.00        | 1.00        | 1.00        | 1.00          | 1.00          | 1.00       | 0.97        |
| <i>blowfish</i> | 0.99       | 1.00        | 1.00        | 1.00        | 0.98        | 1.00          | 0.98          | 1.00       | 0.99        |
|                 | <b>WNS</b> | <b>LUTs</b> | <b>Pins</b> | <b>BRAM</b> | <b>Regs</b> | <b>Blocks</b> | <b>Cycles</b> | <b>DSP</b> | <b>FMax</b> |

**Figure 2:** Relative improvement for all metrics in the *Balanced* scenario

## RESULTS: THE AREA SCENARIO

|                 |            |             |             |             |             |               |               |            |             |
|-----------------|------------|-------------|-------------|-------------|-------------|---------------|---------------|------------|-------------|
| <i>aes</i>      | 0.96       | 0.92        | 1.00        | 1.00        | 0.93        | 1.00          | 0.97          | 1.00       | 0.86        |
| <i>adpcm</i>    | 0.83       | 0.78        | 1.11        | 1.00        | 0.96        | 1.00          | 1.04          | 0.63       | 0.52        |
| <i>sha</i>      | 0.96       | 1.00        | 1.06        | 1.00        | 0.84        | 1.00          | 0.83          | 1.00       | 1.08        |
| <i>motion</i>   | 0.97       | 0.94        | 1.00        | 1.00        | 0.97        | 1.00          | 0.92          | 1.00       | 0.95        |
| <i>mips</i>     | 0.91       | 1.00        | 1.06        | 1.00        | 0.89        | 1.00          | 1.00          | 0.72       | 0.84        |
| <i>gsm</i>      | 0.89       | 0.83        | 1.06        | 1.00        | 0.86        | 1.00          | 0.89          | 0.82       | 1.14        |
| <i>dfsint</i>   | 0.94       | 1.00        | 1.06        | 1.00        | 1.00        | 1.00          | 1.02          | 0.76       | 0.80        |
| <i>dfmul</i>    | 0.82       | 1.00        | 1.06        | 1.00        | 0.90        | 1.00          | 1.21          | 0.27       | 0.86        |
| <i>dfdiv</i>    | 0.77       | 1.00        | 1.22        | 0.67        | 1.03        | 0.67          | 1.62          | 0.28       | 0.77        |
| <i>dfadd</i>    | 0.99       | 1.00        | 1.11        | 1.00        | 0.94        | 1.00          | 1.00          | 1.00       | 1.04        |
| <i>blowfish</i> | 0.99       | 1.00        | 1.00        | 1.00        | 0.97        | 1.00          | 0.91          | 1.00       | 1.01        |
|                 | <b>WNS</b> | <b>LUTs</b> | <b>Pins</b> | <b>BRAM</b> | <b>Regs</b> | <b>Blocks</b> | <b>Cycles</b> | <b>DSP</b> | <b>FMax</b> |

**Figure 3:** Relative improvement for all metrics in the *Area* scenario

## RESULTS: THE PERFORMANCE SCENARIO

|                 |            |             |             |             |             |               |               |            |             |
|-----------------|------------|-------------|-------------|-------------|-------------|---------------|---------------|------------|-------------|
| <i>aes</i>      | 0.82       | 0.75        | 1.00        | 1.00        | 0.92        | 1.00          | 1.05          | 1.00       | 0.63        |
| <i>adpcm</i>    | 0.84       | 0.94        | 1.17        | 1.00        | 0.96        | 1.00          | 0.94          | 0.69       | 0.74        |
| <i>sha</i>      | 0.92       | 1.00        | 1.06        | 1.00        | 0.92        | 1.00          | 0.88          | 1.00       | 0.96        |
| <i>motion</i>   | 0.99       | 1.00        | 1.00        | 1.00        | 1.01        | 1.00          | 1.00          | 1.00       | 0.98        |
| <i>mips</i>     | 0.90       | 1.00        | 1.06        | 1.00        | 0.93        | 1.00          | 1.03          | 0.83       | 0.80        |
| <i>gsm</i>      | 0.95       | 0.92        | 1.00        | 1.00        | 0.95        | 1.00          | 0.90          | 1.00       | 1.02        |
| <i>dfsint</i>   | 0.87       | 1.11        | 1.06        | 1.00        | 1.27        | 1.00          | 1.25          | 0.53       | 0.56        |
| <i>dfmul</i>    | 0.77       | 1.00        | 1.17        | 0.83        | 0.85        | 0.83          | 1.04          | 0.21       | 0.70        |
| <i>dfdiv</i>    | 0.81       | 1.00        | 1.06        | 1.00        | 1.03        | 1.00          | 1.25          | 0.50       | 0.59        |
| <i>dfadd</i>    | 0.97       | 1.00        | 1.00        | 1.00        | 0.97        | 1.00          | 1.00          | 1.00       | 0.94        |
| <i>blowfish</i> | 0.94       | 1.00        | 1.00        | 1.00        | 0.98        | 1.00          | 0.91          | 1.00       | 0.95        |
|                 | <b>WNS</b> | <b>LUTs</b> | <b>Pins</b> | <b>BRAM</b> | <b>Regs</b> | <b>Blocks</b> | <b>Cycles</b> | <b>DSP</b> | <b>FMax</b> |

**Figure 4:** Relative improvement for all metrics in the *Performance* scenario

## RESULTS: THE PERFORMANCE & LATENCY SCENARIO

|                 |            |             |             |             |             |               |               |            |             |
|-----------------|------------|-------------|-------------|-------------|-------------|---------------|---------------|------------|-------------|
| <i>aes</i>      | 0.83       | 0.83        | 1.00        | 1.00        | 0.88        | 1.00          | 0.89          | 1.00       | 0.73        |
| <i>adpcm</i>    | 0.76       | 0.78        | 1.11        | 1.00        | 0.95        | 1.00          | 0.98          | 0.62       | 0.47        |
| <i>sha</i>      | 0.88       | 1.00        | 1.06        | 1.00        | 0.85        | 1.00          | 0.77          | 1.00       | 1.00        |
| <i>motion</i>   | 0.98       | 1.00        | 1.00        | 1.00        | 1.00        | 1.00          | 1.00          | 1.00       | 0.95        |
| <i>mips</i>     | 0.95       | 1.00        | 1.11        | 1.00        | 0.91        | 1.00          | 0.99          | 0.89       | 0.96        |
| <i>gsm</i>      | 0.92       | 0.92        | 1.11        | 1.00        | 0.87        | 1.00          | 0.90          | 0.67       | 1.11        |
| <i>dfsint</i>   | 0.97       | 1.00        | 1.17        | 1.00        | 0.99        | 1.00          | 0.99          | 0.61       | 1.00        |
| <i>dfmul</i>    | 0.86       | 1.00        | 1.22        | 1.00        | 0.89        | 1.00          | 1.01          | 0.33       | 0.84        |
| <i>dfdiv</i>    | 0.82       | 1.00        | 1.11        | 1.00        | 0.84        | 1.00          | 0.96          | 0.34       | 0.83        |
| <i>dfadd</i>    | 0.98       | 1.00        | 1.17        | 1.00        | 0.94        | 1.00          | 0.98          | 1.00       | 1.01        |
| <i>blowfish</i> | 0.96       | 1.00        | 1.00        | 1.00        | 0.97        | 1.00          | 0.93          | 1.00       | 0.97        |
|                 | <b>WNS</b> | <i>LUTs</i> | <i>Pins</i> | <i>BRAM</i> | <i>Regs</i> | <i>Blocks</i> | <i>Cycles</i> | <i>DSP</i> | <i>FMax</i> |

**Figure 5:** Relative improvement for all metrics in the *Performance & Latency* scenario

## RESULTS: COMPARING THE WEIGHTED NORMALIZED SUM (WNS)

|                 |                 |             |                    |                         |
|-----------------|-----------------|-------------|--------------------|-------------------------|
| <i>aes</i>      | 0.94            | 0.96        | 0.82               | 0.83                    |
| <i>adpcm</i>    | 0.84            | 0.83        | 0.84               | 0.76                    |
| <i>sha</i>      | 0.98            | 0.96        | 0.92               | 0.88                    |
| <i>motion</i>   | 0.98            | 0.97        | 0.99               | 0.98                    |
| <i>mips</i>     | 0.95            | 0.91        | 0.90               | 0.95                    |
| <i>gsm</i>      | 0.95            | 0.89        | 0.95               | 0.92                    |
| <i>dfsint</i>   | 0.93            | 0.94        | 0.87               | 0.97                    |
| <i>dfmul</i>    | 0.85            | 0.82        | 0.77               | 0.86                    |
| <i>dfdiv</i>    | 0.84            | 0.77        | 0.81               | 0.82                    |
| <i>dfadd</i>    | 1.00            | 0.99        | 0.97               | 0.98                    |
| <i>blowfish</i> | 0.99            | 0.99        | 0.94               | 0.96                    |
| <b>Average</b>  | 0.93            | 0.91        | 0.89               | 0.90                    |
|                 | <i>Balanced</i> | <i>Area</i> | <i>Performance</i> | <i>Perf. &amp; Lat.</i> |

**Figure 6:** Relative improvement for WNS in all scenarios

In all scenarios:

- Heavily weighted metrics **improved the most**
- Lightly weighted or irrelevant metrics **did not worsen greatly**

### In all scenarios:

- Heavily weighted metrics **improved the most**
- Lightly weighted or irrelevant metrics **did not worsen greatly**

### Autotuning starting point:

- A random start **might achieve more relative improvement**
- But a sensible start **achieves better absolute values**



### Issues with the **Weighted Normalized Sum**:

- The WNS is a **very naive metric combination strategy**
- The **multi-objective optimization** domain provides **better approaches**

### Bigger designs:

- CHStone designs are very simple
- We will use our autotuning approach for bigger FPGA designs

## Bigger designs:

- CHStone designs are very simple
- We will use our autotuning approach for bigger FPGA designs

## Compiler configuration:

- HLS usually includes a C/C++ compiler step
- We will include compiler parameters or passes in our autotuner

## Bigger designs:

- CHStone designs are very simple
- We will use our autotuning approach for bigger FPGA designs

## Compiler configuration:

- HLS usually includes a C/C++ compiler step
- We will include compiler parameters or passes in our autotuner

## Other HLS tools:

- Other HLS tools provide different search spaces
- We will use our autotuning approach for more proprietary and open HLS tools

## Autotuning HLS for FPGAs:

- Designers can **select optimization objectives using weights**
- Good **WNS improvements** after **1.5 hours** of autotuning
- Improvements for targeted metrics **without worsening other metrics**
- A **sensible starting point** impacts autotuning results positively

# AUTOTUNING HLS FOR FPGAS USING OPENTUNER AND LEGUP

---

Pedro Bruel ([phrb@ime.usp.br](mailto:phrb@ime.usp.br))

**Alfredo Goldman** ([gold@ime.usp.br](mailto:gold@ime.usp.br))

Sai Rahul Chalamalasetti ([gold@ime.usp.br](mailto:gold@ime.usp.br))

Dejan Milojicic ([gold@ime.usp.br](mailto:gold@ime.usp.br))

ReConFig, December 5, 2017



*Institute of Mathematics and Statistics  
University of São Paulo*

