

AUTOTUNING HLS FOR FPGAS USING OPENTUNER AND LEGUP

Pedro Bruel (phrb@ime.usp.br)

Alfredo Goldman (gold@ime.usp.br)

Sai Rahul Chalamalasetti (gold@ime.usp.br)

Dejan Milojicic (gold@ime.usp.br)

ReConFig, December 5, 2017



*Institute of Mathematics and Statistics
University of São Paulo*



1. FPGAs, HLS & Autotuning
2. Background
3. Experiments & Results
4. Conclusion



The slides and all source code are hosted at [GitHub](#):

- [Code & Data](#): `github.com/phrb/legup-tuner`
- [Slides](#): `github.com/phrb/slides-reconfig-2017-autotuning`

FPGAs:

- Logic Blocks and Interconnections
- Reconfigurable

Tradeoff:

- Energy Efficiency and Performance
- Programmability



Using FPGAs in Bing:

1,632 Servers with FPGAs Running Bing Page Ranking Service (~30,000 lines of C++)

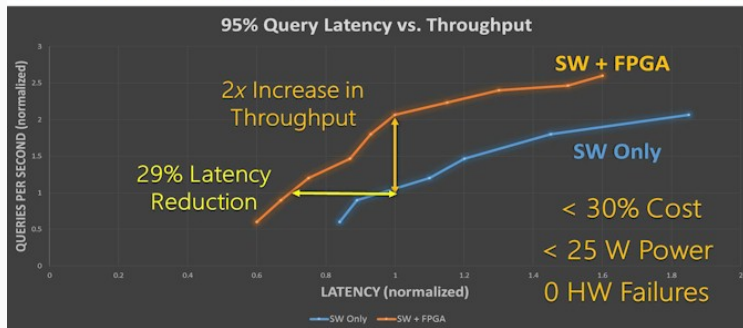


Image: enterprisetech.com/2014/09/03/microsoft-using-fpgas-speed-bing-search/ [Accessed in 27/11/17]

FPGAs: HIGH-LEVEL SYNTHESIS

LegUp HLS flow:




Image: Canis, Andrew Christopher. LegUp: Open-Source High-Level Synthesis Research Framework. Diss. University of Toronto, 2015.

FPGAs: HIGH-LEVEL SYNTHESIS

HLS can generate **lower-latency applications**:




Image: Karras, Kimon, Michaela Blott, and Kees Vissers. "High-Level Synthesis Case Study: Implementation of a Memcached Server." arXiv preprint arXiv:1408.5387 (2014).

FPGAs: HIGH-LEVEL SYNTHESIS

Qualitatively, with less effort:

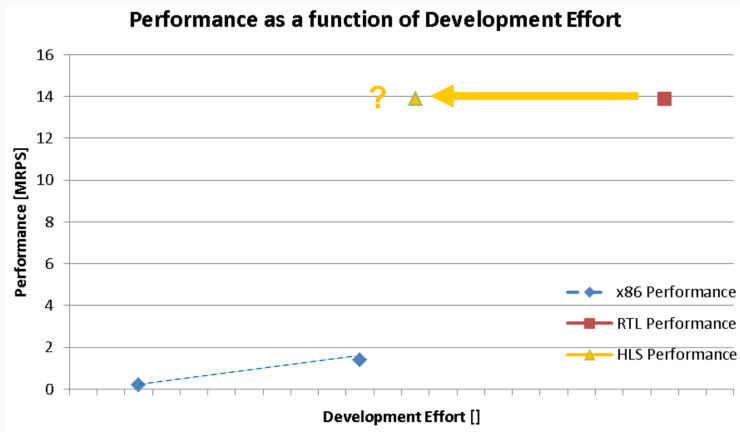


Image: Blott, Michaela, et al. "Achieving 10Gbps line-rate key-value stores with FPGAs." Presented as part of the 5th USENIX Workshop on Hot Topics in Cloud Computing. 2013.

This is an **old issue**:




Image: Smith, Steven W. "The scientist and engineer's guide to digital signal processing." 1997

Why use autotuning for HLS?

Compare with Huang's work

Describe Xu's work with OpenTuner

Metric composition problem:

- 8 hardware metrics obtained from Quartus
- But the autotuner optimizes a single value

Metric composition problem:

- 8 hardware metrics obtained from Quartus
- But the autotuner optimizes a single value

Our solution so far:

- Normalize metrics using initial values
- Optimize the normalized sum of the 8 metrics

Normalized sum of metrics:

$$f(M, W) = \frac{\sum_{\substack{m_i \in M \\ w_i \in W}} w_i \left(\frac{m_i}{m_i^0} \right)}{\sum_{w_i \in W} w_i}$$

Where M is the set of **hardware metrics** and W is the set of **weights**

An **Optimization Scenario** consists of:

- An **optimization objective**: performance, area, . . .
- **Weights** for hardware metrics

An **Optimization Scenario** consists of:

- An **optimization objective**: performance, area, . . .
- **Weights** for hardware metrics

Our scenarios:

- 3 **specific** scenarios & 1 **balanced** scenario
- Weights: **powers of two** from 1: **irrelevant** to 8: **high**

WEIGHTED OPTIMIZATION SCENARIOS

Table 1: [Weights](#) for [Optimization Scenarios](#) ([High](#) = 8, [Medium](#) = 4, [Low](#) = 2)

Metric	Area	Perf. & Lat	Performance	Balanced
<i>LUT</i>	High	Low	Low	Medium
<i>Registers</i>	High	High	Medium	Medium
<i>BRAMs</i>	High	Low	Low	Medium
<i>DSPs</i>	High	Low	Low	Medium
<i>FMax</i>	Low	High	High	Medium
<i>Cycles</i>	Low	High	Low	Medium

Experiments

- We performed 10 autotuning runs for each scenario
- Each run took 1.5 hours
- We kept track of how each hardware metric changed over time
- We used the Weighted Normalized Sum (WNS) as autotuning objective

Presentation of results for each scenario:

- Mean of the relative changes, over 10 tuning runs
- For each hardware metric, and for WNS
- For each CHStone application
- Bluer is always better than Redder

RESULTS: COMPARING DEFAULT & RANDOM STARTS

<i>aes</i>	0.79	0.91	1.00	0.56	1.00	0.47	1.00	1.12
<i>adpcm</i>	0.68	1.13	1.00	0.54	1.00	0.56	0.60	0.98
<i>sha</i>	1.00	1.03	1.00	0.82	1.00	0.55	1.00	0.89
<i>motion</i>	1.02	1.00	0.60	0.85	1.00	0.57	1.00	0.94
<i>mips</i>	1.00	0.93	1.00	0.44	1.00	0.45	0.98	1.24
<i>gsm</i>	0.83	1.17	1.00	0.48	1.00	0.56	0.52	0.99
<i>dfsint</i>	0.79	0.97	1.00	0.61	1.00	0.60	1.49	1.53
<i>dfmul</i>	1.00	1.06	0.90	0.47	0.90	0.47	1.31	1.10
<i>dfdiv</i>	0.83	1.07	0.80	0.73	0.80	0.65	1.32	1.49
<i>dfadd</i>	1.00	0.94	1.00	0.82	1.00	0.71	1.00	0.93
<i>blowfish</i>	—	—	—	—	—	—	—	—
	<i>LUTs</i>	<i>Pins</i>	<i>BRAM</i>	<i>Regs</i>	<i>Blocks</i>	<i>Cycles</i>	<i>DSP</i>	<i>FMax</i>

Figure 1: Comparison of the absolute values for Random and Default starting points in the Balanced scenario

RESULTS: THE BALANCED SCENARIO

<i>aes</i>	0.94	0.90	1.00	1.00	0.97	1.00	0.98	1.00	0.76
<i>adpcm</i>	0.84	0.73	1.13	1.00	0.91	1.00	1.05	0.63	0.49
<i>sha</i>	0.98	1.00	1.03	1.00	0.96	1.00	0.86	1.00	0.97
<i>motion</i>	0.98	0.97	1.00	1.00	0.99	1.00	0.95	1.00	0.95
<i>mips</i>	0.95	1.00	1.07	1.00	0.90	1.00	1.01	0.80	0.89
<i>gsm</i>	0.95	0.95	1.17	1.00	0.99	1.00	0.95	0.49	1.08
<i>dfsint</i>	0.93	1.03	1.03	1.00	1.05	1.00	1.07	0.65	0.73
<i>dfmul</i>	0.85	1.00	1.13	0.90	0.88	0.90	1.06	0.31	0.76
<i>dfdiv</i>	0.84	1.00	1.07	0.80	1.15	0.80	1.34	0.41	0.57
<i>dfadd</i>	1.00	1.00	1.00	1.00	1.00	1.00	1.00	1.00	0.97
<i>blowfish</i>	0.99	1.00	1.00	1.00	0.98	1.00	0.98	1.00	0.99
	WNS	LUTs	Pins	BRAM	Regs	Blocks	Cycles	DSP	FMax

Figure 2: Relative improvement for all metrics in the *Balanced* scenario

RESULTS: THE AREA SCENARIO

<i>aes</i>	0.96	0.92	1.00	1.00	0.93	1.00	0.97	1.00	0.86
<i>adpcm</i>	0.83	0.78	1.11	1.00	0.96	1.00	1.04	0.63	0.52
<i>sha</i>	0.96	1.00	1.06	1.00	0.84	1.00	0.83	1.00	1.08
<i>motion</i>	0.97	0.94	1.00	1.00	0.97	1.00	0.92	1.00	0.95
<i>mips</i>	0.91	1.00	1.06	1.00	0.89	1.00	1.00	0.72	0.84
<i>gsm</i>	0.89	0.83	1.06	1.00	0.86	1.00	0.89	0.82	1.14
<i>dfsint</i>	0.94	1.00	1.06	1.00	1.00	1.00	1.02	0.76	0.80
<i>dfmul</i>	0.82	1.00	1.06	1.00	0.90	1.00	1.21	0.27	0.86
<i>dfdiv</i>	0.77	1.00	1.22	0.67	1.03	0.67	1.62	0.28	0.77
<i>dfadd</i>	0.99	1.00	1.11	1.00	0.94	1.00	1.00	1.00	1.04
<i>blowfish</i>	0.99	1.00	1.00	1.00	0.97	1.00	0.91	1.00	1.01
	WNS	LUTs	Pins	BRAM	Regs	Blocks	Cycles	DSP	FMax

Figure 3: Relative improvement for all metrics in the *Area* scenario

RESULTS: THE PERFORMANCE SCENARIO

<i>aes</i>	0.82	0.75	1.00	1.00	0.92	1.00	1.05	1.00	0.63
<i>adpcm</i>	0.84	0.94	1.17	1.00	0.96	1.00	0.94	0.69	0.74
<i>sha</i>	0.92	1.00	1.06	1.00	0.92	1.00	0.88	1.00	0.96
<i>motion</i>	0.99	1.00	1.00	1.00	1.01	1.00	1.00	1.00	0.98
<i>mips</i>	0.90	1.00	1.06	1.00	0.93	1.00	1.03	0.83	0.80
<i>gsm</i>	0.95	0.92	1.00	1.00	0.95	1.00	0.90	1.00	1.02
<i>dfsint</i>	0.87	1.11	1.06	1.00	1.27	1.00	1.25	0.53	0.56
<i>dfmul</i>	0.77	1.00	1.17	0.83	0.85	0.83	1.04	0.21	0.70
<i>dfdiv</i>	0.81	1.00	1.06	1.00	1.03	1.00	1.25	0.50	0.59
<i>dfadd</i>	0.97	1.00	1.00	1.00	0.97	1.00	1.00	1.00	0.94
<i>blowfish</i>	0.94	1.00	1.00	1.00	0.98	1.00	0.91	1.00	0.95
	WNS	LUTs	Pins	BRAM	Regs	Blocks	Cycles	DSP	FMax

Figure 4: Relative improvement for all metrics in the *Performance* scenario

RESULTS: THE PERFORMANCE & LATENCY SCENARIO

<i>aes</i>	0.83	0.83	1.00	1.00	0.88	1.00	0.89	1.00	0.73
<i>adpcm</i>	0.76	0.78	1.11	1.00	0.95	1.00	0.98	0.62	0.47
<i>sha</i>	0.88	1.00	1.06	1.00	0.85	1.00	0.77	1.00	1.00
<i>motion</i>	0.98	1.00	1.00	1.00	1.00	1.00	1.00	1.00	0.95
<i>mips</i>	0.95	1.00	1.11	1.00	0.91	1.00	0.99	0.89	0.96
<i>gsm</i>	0.92	0.92	1.11	1.00	0.87	1.00	0.90	0.67	1.11
<i>dfsint</i>	0.97	1.00	1.17	1.00	0.99	1.00	0.99	0.61	1.00
<i>dfmul</i>	0.86	1.00	1.22	1.00	0.89	1.00	1.01	0.33	0.84
<i>dfdiv</i>	0.82	1.00	1.11	1.00	0.84	1.00	0.96	0.34	0.83
<i>dfadd</i>	0.98	1.00	1.17	1.00	0.94	1.00	0.98	1.00	1.01
<i>blowfish</i>	0.96	1.00	1.00	1.00	0.97	1.00	0.93	1.00	0.97
	WNS	LUTs	Pins	BRAM	Regs	Blocks	Cycles	DSP	FMax

Figure 5: Relative improvement for all metrics in the *Performance & Latency* scenario

RESULTS: COMPARING THE WEIGHTED NORMALIZED SUM (WNS)

<i>aes</i>	0.94	0.96	0.82	0.83
<i>adpcm</i>	0.84	0.83	0.84	0.76
<i>sha</i>	0.98	0.96	0.92	0.88
<i>motion</i>	0.98	0.97	0.99	0.98
<i>mips</i>	0.95	0.91	0.90	0.95
<i>gsm</i>	0.95	0.89	0.95	0.92
<i>dfsint</i>	0.93	0.94	0.87	0.97
<i>dfmul</i>	0.85	0.82	0.77	0.86
<i>dfdiv</i>	0.84	0.77	0.81	0.82
<i>dfadd</i>	1.00	0.99	0.97	0.98
<i>blowfish</i>	0.99	0.99	0.94	0.96
Average	0.93	0.91	0.89	0.90
	<i>Balanced</i>	<i>Area</i>	<i>Performance</i>	<i>Perf. & Lat.</i>

Figure 6: Relative improvement for WNS in all scenarios

In all scenarios:

- Heavily weighted metrics improved the most
- Lightly weighted or irrelevant metrics did not worsen greatly

In all scenarios:

- Heavily weighted metrics improved the most
- Lightly weighted or irrelevant metrics did not worsen greatly

Autotuning starting point:

- A random start might achieve more relative improvement
- But a sensible start achieves better absolute values

LIMITATIONS OF THIS WORK

Discuss the issues with the weighted cost function

Discuss all future work topics

AUTOTUNING HLS FOR FPGAS USING OPENTUNER AND LEGUP

Pedro Bruel (phrb@ime.usp.br)

Alfredo Goldman (gold@ime.usp.br)

Sai Rahul Chalamalasetti (gold@ime.usp.br)

Dejan Milojicic (gold@ime.usp.br)

ReConFig, December 5, 2017



*Institute of Mathematics and Statistics
University of São Paulo*

